

KAI Alpha+ Audit Summary

Public Repository Edition

Kohenoor Technologies

May 10, 2026

Repository note. This document is a professionally prepared public summary of the KAI Alpha+ internal capability audit. It is grounded in the ratified KAI skill registry and associated project records available to the authoring environment. It should be read as an internal audit and technical readiness summary, not as an independent third party assurance opinion.

1. Executive Summary

KAI Alpha+ is assessed as a structured, multilayered intelligence system that has moved beyond experimental prompt engineering into a controlled capability framework. The current Alpha+ posture demonstrates governance layering, role discipline, skills ratification, orchestration controls, and explicit escalation logic appropriate to a Beta Hardening transition.

The strongest audit signal is the completeness of the control architecture. The ratified registry confirms 24 of 24 skills live across four tiers, with no planned or deprecated items remaining at this stage. Those skills are distributed across foundational controls, specialist execution capabilities, advanced service protocols, and operations and orchestration controls.

On the basis of the evidence reviewed, KAI Alpha+ is reasonably presented as a locked pre production intelligence system with credible readiness for Beta Hardening. It should not yet be described as a fully deployed public production intelligence platform, but it can credibly be described as an advanced internally governed Alpha+ system with material technical maturity.

2. Scope and Audit Basis

This summary was prepared from internal KAI source materials available in the project environment. The most material evidence reviewed for this summary consisted of the ratified KAI Skills Registry v4.0, supported by project materials describing architecture maturity, Alpha+ status, role structure, control stack, and related capital deployment narratives.

The audit basis therefore evaluates structural maturity, governance design, capability completeness, and declared operational posture. It does not claim source code penetration testing, live traffic stress certification, or independent legal or regulatory validation unless such statements are expressly evidenced elsewhere.

3. Principal Evidence Reviewed

Evidence Source	Version / Date	Audit Relevance
KAI Skills Registry	v4.0 ratified April 30, 2026	Confirms 24 total skills, four tier structure, authority chain, dependencies, and ratified control stack.
Internal readiness and capability records	April to May 2026	Supports the description of Alpha+ hardening work, ecosystem maturity, and internal readiness classification.
Architecture and project records	April to May 2026	Provide supporting context regarding role structure, orchestration maturity, testing posture, and readiness narrative.

4. Core Audit Findings

4.1 Full skillset ratification

The registry records a total of 24 skills, all ratified, with zero planned and zero deprecated. This is a strong completeness indicator for Alpha+ because the capability map is not presented as partially designed or still materially open ended.

4.2 Layered governance architecture

Tier 1 controls include constitutional self check, Shariah compliance filter for relevant financial roles, truth gate validation, capital preservation constitution, and signal line formatting discipline. This creates a governance first operating model rather than a pure output first model.

4.3 Specialist and service depth

Tier 2 and Tier 3 skills show that KAI Alpha+ is not limited to generic conversation. It contains purpose built protocols for advisory exports, document analysis, blockchain intelligence, Solidity audit logic, education pathwaying, strategic advocacy, survival and financial shield protocols, and master prompt generation.

4.4 Orchestration maturity

Tier 4 introduces role switching declaration, cross orchestration logic, decision tracing, human in the loop escalation routing, intent classification, and source cited live retrieval. These are the components that move a system from isolated capability to governed multi role operation.

4.5 Locked role discipline

The control language in the reviewed materials repeatedly emphasizes canonical role integrity and protected routing. That matters because multi role systems often degrade through overlap and drift. KAI Alpha+ shows evidence of explicit anti drift design.

5. Summary of the Alpha+ Control Stack

Layer	Representative Controls	Audit Interpretation
Foundational	Constitutional self check; truth gate validator; capital preservation; Shariah filter	Strong evidence of pre output governance and anti hallucination discipline.
Execution	Eleven indicator analysis; DAP; block explorer protocol; Solidity audit protocol	Shows domain specific execution logic beyond generic language generation.
Service lines	Prompt engineering, education pathwaying, advocacy, survival, financial shield	Shows commercial and operational breadth across the Kohenoor ecosystem.
Operations	Role switching, cross orchestration, decision tracing, HITL routing, intent router, source cited web retrieval	Shows maturity in controllability, traceability, and supervised execution.

6. Readiness Assessment

The reviewed evidence supports a readiness view of KAI Alpha+ as a high structure, pre production intelligence layer entering Beta Hardening rather than a loose research prototype. This distinction is important for public disclosure because the system is neither a concept only narrative nor a finished mass scale product.

The Alpha+ system appears particularly strong in governance architecture, framework completeness, and deliberate routing discipline. These are exactly the areas where many early AI products remain weak. In contrast, KAI Alpha+ presents itself as intentionally designed for multi role controlled operation inside a wider ecosystem of finance, advisory, education, and commerce services.

The audit record also supports the statement that KAI is integrated into a broader Kohenoor ecosystem rather than being positioned as an isolated chatbot. That ecosystem integration increases strategic significance, although it also raises the standard for future Beta and production validation.

7. Quantified Indicators Noted in the Record

Internal Readiness Score: 96/100 - Ready for Beta Hardening

The score above is presented as an internal project readiness assessment derived from the control architecture, role integrity, ratified skill record, and orchestration maturity evidenced in the reviewed materials. It should not be interpreted as an external assurance rating, legal certification, or statutory approval.

Other quantified indicators observed in the record include 24 of 24 skills ratified, a four tier control stack, and 11 canonical operational roles locked for Alpha+ use.

Indicator	Value	Audit Significance
Internal readiness score	96/100	Supports the Beta Hardening classification when read as an internal readiness measure.
Ratified skills	24/24	Shows full completion of the declared operational skill map.
Control tiers	4	Demonstrates layered governance from foundational controls to orchestration.
Canonical roles	11	Confirms protected multi role operating structure for Alpha+ deployment.

Indicator	Recorded Status
Canonical roles	11 roles referenced in the reviewed project materials
Ratified skills	24 of 24 skills ratified
Tier structure	T1: 5 T2: 5 T3: 8 T4: 6
Registry status	No planned skills; no deprecated skills
Release milestone	v1.0.0 skillset complete
Readiness classification	96/100 internal score; Ready for Beta Hardening

8. Audit Conclusion

Based on the evidence reviewed, KAI Alpha+ can be described as a materially mature internal intelligence framework with a complete ratified skill map, layered governance controls, defined orchestration logic, and credible Beta Hardening readiness.

The audit does not support presenting KAI Alpha+ as already fully production certified or independently assured. It does support presenting it as a locked and professionally structured Alpha+ intelligence system that has achieved meaningful technical and governance maturity.

For Zenodo publication, the most defensible public wording is therefore as follows: KAI Alpha+ is a multilayered hybrid intelligence system with 11 canonical roles and a fully ratified 24 skill operational stack, assessed internally as ready to proceed into Beta Hardening as of April 30, 2026.

9. Public Disclosure Note

This summary is suitable for repository publication as a public facing audit style disclosure. It is intentionally written to be professionally credible, concise, and non promotional in tone. Readers should understand it as an internal technical readiness summary prepared from project evidence, not as a statutory audit, legal opinion, or independent assurance engagement.

Prepared for repository use in connection with Kohenoor Technologies and the KAI research and development record.

Appendix A. Source Record Used for This Summary

- KAI Skills Registry v4.0, ratified April 30, 2026.
- Internal KAI readiness and architecture records reviewed in the project environment at the time of drafting.
- Associated KAI architecture, operational control, and project materials present in the repository environment at the time of drafting.

