

KAI Multilayered LLM Orchestration

Repository Note for Research Record and Product Evidence

Version	Phase	Date	Prepared For
v1.0	Alpha+ / Pre Beta Hardening	28 May 2026	KAI Research Repository

1. Purpose of this Repository Note

This note records the KAI user interface evidence showing **Multilayered LLM Orchestration** as part of the Kohenoor AI research and product development repository. The captured screen confirms that KAI presents itself as an orchestrator that can detect user intent, activate the right specialist role, and route a query to a suitable LLM engine without manual role selection.

2. Captured Interface Evidence

The following screenshot should be stored with the repository entry as visual evidence of the current Alpha+ interface state.




Figure 1: KAI interface showing the “Simple” mode response about LLM orchestration families and complexity tiers.

3. Description

KAI is presented as a **Multilayered LLM Orchestrator** within the Kohenoor ecosystem. KAI can intelligently detect whether the user needs legal analysis, code support, market analysis, compliance support, blockchain analysis, or general assistance, and then route the query to the appropriate specialist role and LLM family.

4. LLM Families Shown in the Captured Response

LLM Family	Model / Variant	Primary Use in KAI	Complexity Tier
DeepSeek	deepseek-chat, deepseek-reasoner	Default engine for general queries, coding, and fast responses.	Tier 1 and Tier 2
Qwen by Alibaba	qwen-plus, qwen-coder-plus	Moderate complexity general and coding tasks.	Tier 2 and Tier 3
Kimi by Moonshot AI	moonshot-v1-32k, moonshot-v1-128k	Long context research and strategy.	Tier 3
Gemini 1.5 Pro by Google	Gemini 1.5 Pro	Complex market intelligence, financial analysis, and long documents.	Tier 3
GPT-4o by OpenAI	GPT-4o	High complexity general, financial, and market queries.	Tier 4
Claude Opus by Anthropic	Claude Opus	Highest stakes legal, compliance, blockchain, and survival analysis.	Tier 4

5. Complexity Tier Logic Captured

The interface states that complexity tiers are assigned automatically: **T1** for simple or short queries, **T2** for moderate queries, **T3** for complex or long queries, and **T4** for high stakes domains. This supports the repository claim that KAI's orchestration is not merely a manual model selector, but an intent and complexity based routing layer.

6. Research Significance

- Documents KAI's transition from single model interaction to multi model orchestration.
- Supports the Alpha+ evidence trail for KAI Nexus and KAI Super Model development.
- Shows a public safe description of role based and LLM based routing without disclosing proprietary prompts, keys, internal scoring, or routing logic.
- Provides a useful reference for future Beta Hardening audits, investor briefings, and repository disclosures.

7. Repository Abstract

This repository entry records the Alpha+ interface evidence of KAI's Multilayered LLM Orchestration. KAI is designed to detect user intent, assign the request to the appropriate specialist role, and route the query through a suitable LLM family based on complexity and domain sensitivity. The current interface identifies six LLM families used during Alpha+ testing: DeepSeek, Qwen, Kimi, Gemini, GPT-4o, and Claude Opus. The captured evidence supports the ongoing Beta Hardening narrative that KAI is evolving into a role aware, model aware, and complexity aware AI intelligence layer for the Kohenoor ecosystem.